

Microservices

Where Did It All Go Wrong

Ian Cooper @ICooper (X, Hachyderm.io)

Who are you?

I am a polyglot coding architect with over 20 years of experience delivering solutions in government, healthcare, and finance and ecommerce. During that time I have worked for the DTI, Reuters, Sungard, Misys, Beazley, Huddle and Just Eat Takeaway delivering everything from bespoke enterprise solutions, 'shrink-wrapped' products for thousands of customers, to SaaS applications for hundreds of thousands of customers.

I am an experienced systems architect with a strong knowledge of OO, TDD/BDD, DDD, EDA, CQRS/ES, REST, Messaging, Design Patterns, Architectural Styles, ATAM, and Agile Engineering Practices

I am frequent contributor to OSS, and I am the owner of: <https://github.com/BrighterCommand>. I speak regularly at user groups and conferences around the world on architecture and software craftsmanship. I run public workshops teaching messaging, event-driven and reactive architectures.

I have a strong background in C#. I spent years in the C++ trenches. I dabble in Go, Java, JavaScript and Python.

www.linkedin.com/in/ian-cooper-2b059b



B R I G H T E R

Agenda

- Fundamentals & Definitions
- The Problem
- Advice
- Q&A

Fundamentals & Definitions

When I use a word,' Humpty Dumpty said in rather a scornful tone, 'it means just what I choose it to mean — neither more nor less.

— Lewis Carroll, Through the Looking Glass

Fundamental Theorem of Software Engineering

$$C(P) > \left(\frac{1}{2}P\right) + C\left(\frac{1}{2}P\right)$$

Due to complexity: It is always easier (and cheaper) to create two small pieces rather than one big piece if the two small pieces do the same job as the single piece. This is because combining two problems in one piece forces us to deal with the unneeded interactions between the two pieces

Fundamental Theorem of Software Engineering

Specifically, we suggest that the cost of implementing a computer system will be minimized when the parts of the problem are:

- solvable separately
- manageably small

Of course, everyone has a different definition of "manageably small"

module: A module is a lexically contiguous sequence of program statements, bounded by boundary elements, having an aggregate identifier. Another way of saying this is that a module is a bounded, contiguous group of statements having a single name by which it can be referred to as a unit.

software component: software components are things that are independently replaceable and upgradeable

- Martin Fowler, [Software Component](#)

- components come in two forms: libraries and services.
- components relate to modules, but are replaceable

service: A service is a component that exists in its own process, clients talk to it over some inter-process communication mechanism: RPC, RESTful calls over HTTP, messaging, etc.

- Martin Fowler, [Software Component](#)

For services to be components:

- The upgrade/replacement of one service should not require co-ordination with another

microservices: a set of practices for successfully building a *software component* architectural style (at scale):

- Componentization via Services*
- Organized around Business Capabilities
- Smart endpoints and dumb pipes
- Design for failure
- Infrastructure Automation
- Decentralized Data Management
- Decentralized Governance
- Products not Projects

Lewis, James Fowler ,Martin - [Microservices](#)

* It is a component and a service i.e. should be independently replaceable (deployable)

Benefits

- Strong Module Boundaries
- Independent Deployment
- Technology Diversity
- Fault Tolerance
- Cognitive Load (Code that fits in my head)

Costs

- Distribution
- Verification
- Operational Complexity
- Skill Issues
- Data: BI & ML
- After Fowler, Martin [Microservice Trade-Offs](#)

Microservice Premium

don't even consider microservices unless you have a system that's too complex to manage as a monolith

- Fowler, Martin [Microservice Premium](#)

- Code doesn't fit in your head
- At least a Five-Pizza Team
- You understand your sub-domains
- Sub-domains have different forces for change
- A wide range of interactions

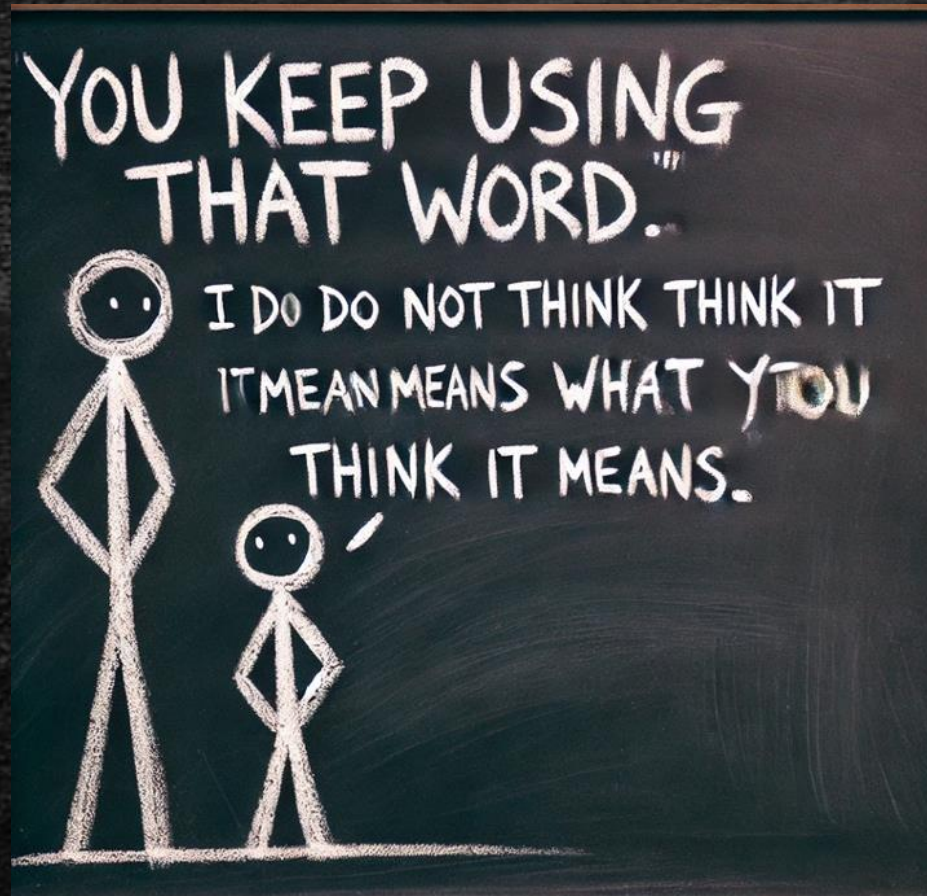
You are solving ****large**** problems

The Problem

Parents wonder why the streams are bitter, when they themselves poison the fountain.

— John Locke,


Microservices



- "They are very, very small. I mean 100 lines of code is probably a big service these days." Fred George, Barcelona Ruby Conference
- "If its more than one programmer to develop & design and maintain it, it's not a microservice" Fred George, GOTO 2016
- "Single **function** microservices provide us with the highest degree of granularity. This increase in granularity gives us more control over an application's resource consumption, availability, and scalability." – Think Microservices

- We wanted REPLACEABLE not SMALL. SIZE is just a factor in replaceability!!!


Nanoservices



Nanoservice – Where the overhead of a being service not balanced by the value.

- Intramodule effects
- Intermodule effects
- Total Cost

Edward Yourdon, Larry Constantine,
Structured Design 1975



Fracture Plane: Team Availability

A fracture plane is a natural seam in the software system that allows the system to be split easily into two or more parts....

Skelton, Matthew; Pais, Manuel. Team Topologies

Fracture plane: team availability is an anti-pattern that occurs when the driving decision is that we want to create a fracture plane to support a new team working on the problem over assigning to an existing team, for the whom the work may not represent the highest priority, to achieve faster time-to-market.

Cost of Ownership

The cost of ownership of software exceeds the cost of authorship (proof: assume a two-pizza team takes 3 months to write software but 36 months owning it, the cost of ownership is an order of magnitude greater).

Reducing the cost of authorship improves the time to market. In a context where speed of feature release wins we risk prioritizing fracture plane: team availability.

But over time a growing number of components will bankrupt us as ownership costs rise-and rise.

Microservices




Microservice

=



Process

Each service is autonomous and self-contained and runs a unique process. [PaloAltoNetworks.com](https://www.paloaltonetworks.com)



Microservice

=



Container

Each **microservice** is packaged as a **Docker container** to enable deployment to a **Kubernetes** cluster for application orchestration. [ThinkMicroservices.com](https://www.thinkmicroservices.com)



Microservice

=



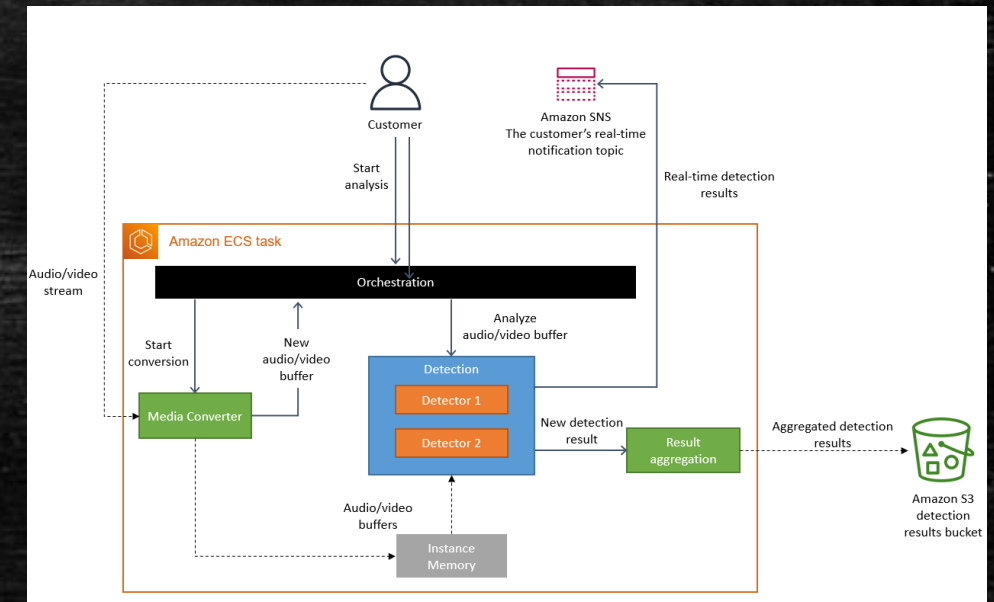
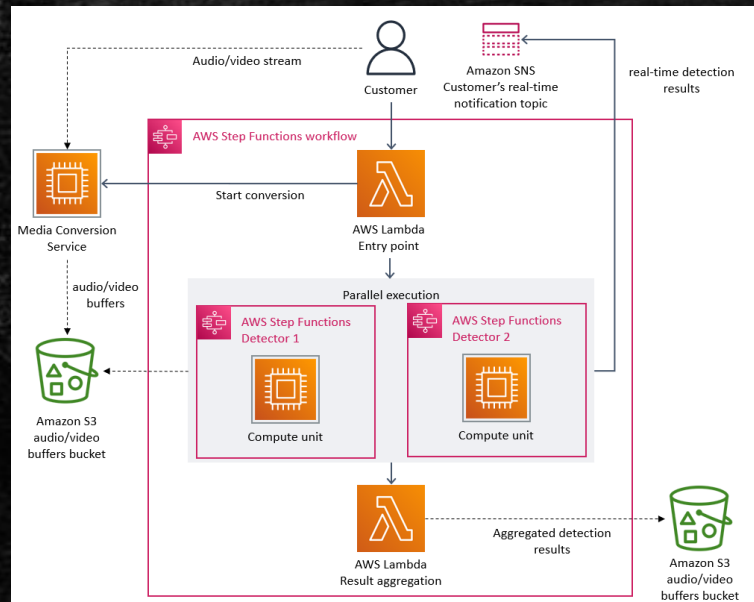
FaaS

functions .. launched by AWS Lambda, is all that you need to build a microservice ... to the level of granularity desired [AWS.Amazon.com](https://aws.amazon.com)

Prime Video

From distributed microservices to a monolith application

We realized that distributed approach wasn't bringing a lot of benefits in our specific use case, so we packed all of the components into a single process.




Even Amazon can't make sense of serverless or microservices

That really sums up so much of the microservices craze that was tearing through the tech industry for a while: IN THEORY. Now the real-world results of all this theory are finally in, and **it's clear that in practice, microservices pose perhaps the biggest siren song for needlessly complicating your system.** And serverless only makes it worse.

— Hansson, David Heinemeier [blog](#)

YOU MUST BE THIS TALL
TO USE MICROSERVICES.




Advice

Why is it that a child sometimes does the opposite of what he is told? Why would a person sometimes dislike receiving a favor? Why is propaganda frequently ineffective in persuading people? And why would the grass in the adjacent pasture ever appear greener?.

— Jack Brehm, 1966

The “4 +1 “ View Model - Philippe Kruchten

t: 18



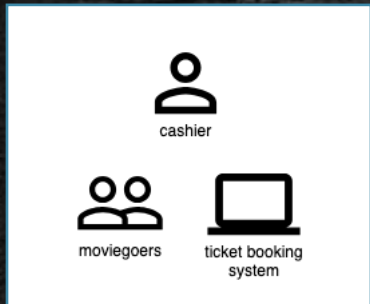
Scenarios

The elements in the four views are shown to work together seamlessly by the use of a small set of important scenarios —instances of more general *use cases*... as a driver to discover the architectural elements during the architecture design as we will describe later

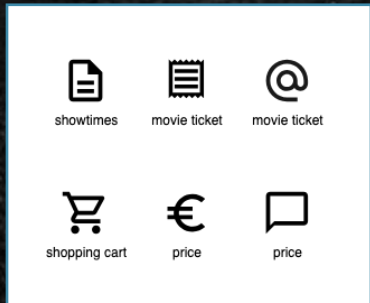
Kruchten, Philippe Architectural Blueprints the 4 +1 View Model of Software

This is where we capture our requirements


Scenarios (Functional Requirements)



Actors: systems or people the story is about



Work Objects: actors work with or exchange information about work objects: resources in the domain



Activities: An activity describes the relationship of the actor to the work object – what do they do to it?

Domain Storytelling

Scenarios (Quality Attributes)


Table 1: Structure of Matrix to Evaluate Candidate Patterns

	Pattern 1		Pattern 2		...	Pattern n	
	Pros	Cons	Pros	Cons		Pros	Cons
Architectural driver 1							
Architectural driver 2							
...							
Architectural driver n							

Attribute Driven Design

The “4 +1 “ View Model - Philippe Kruchten

t: 20



- See Chris Simon, [Experiences scaling a modular monolith to microservices using the 4+1 views](#)

Logical

The logical architecture primarily supports the functional requirements—what the system should provide in terms of services to its users. The system is decomposed into a set of key abstractions, taken (mostly) from the problem domain

Kruchten, Philippe Architectural Blueprints the 4 +1 View Model of Software

This is where we explore our requirements, typically a domain model

Logical

A fracture plane is a natural seam in the software system that allows the system to be split easily into two or more parts.... It is usually best to try to align software boundaries with the different business domain areas...

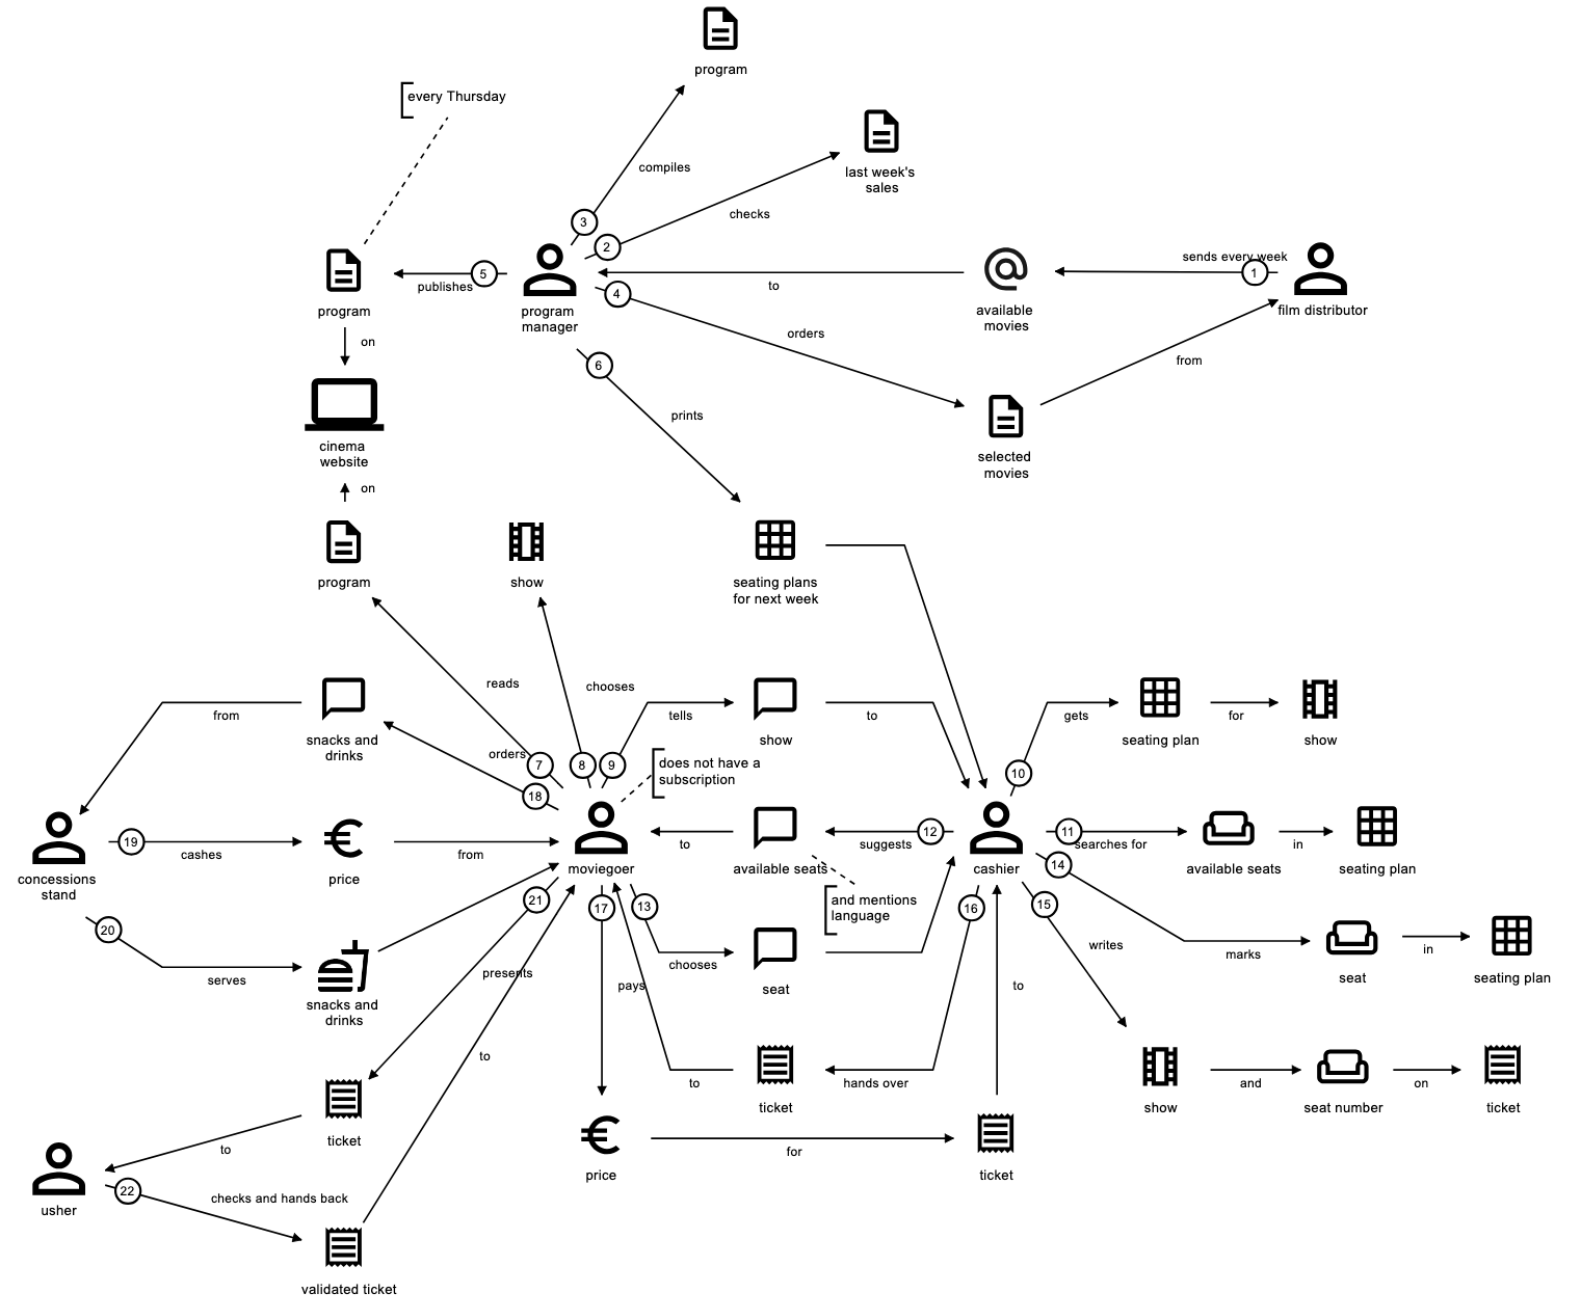
Skelton, Matthew; Pais, Manuel. Team Topologies

How do we find these sub-domains within our system?

- Natural seams within the domain understood by SMEs
- Domain Storytelling
- Value Stream Mapping to align with processes

Domain Storytelling

Examples



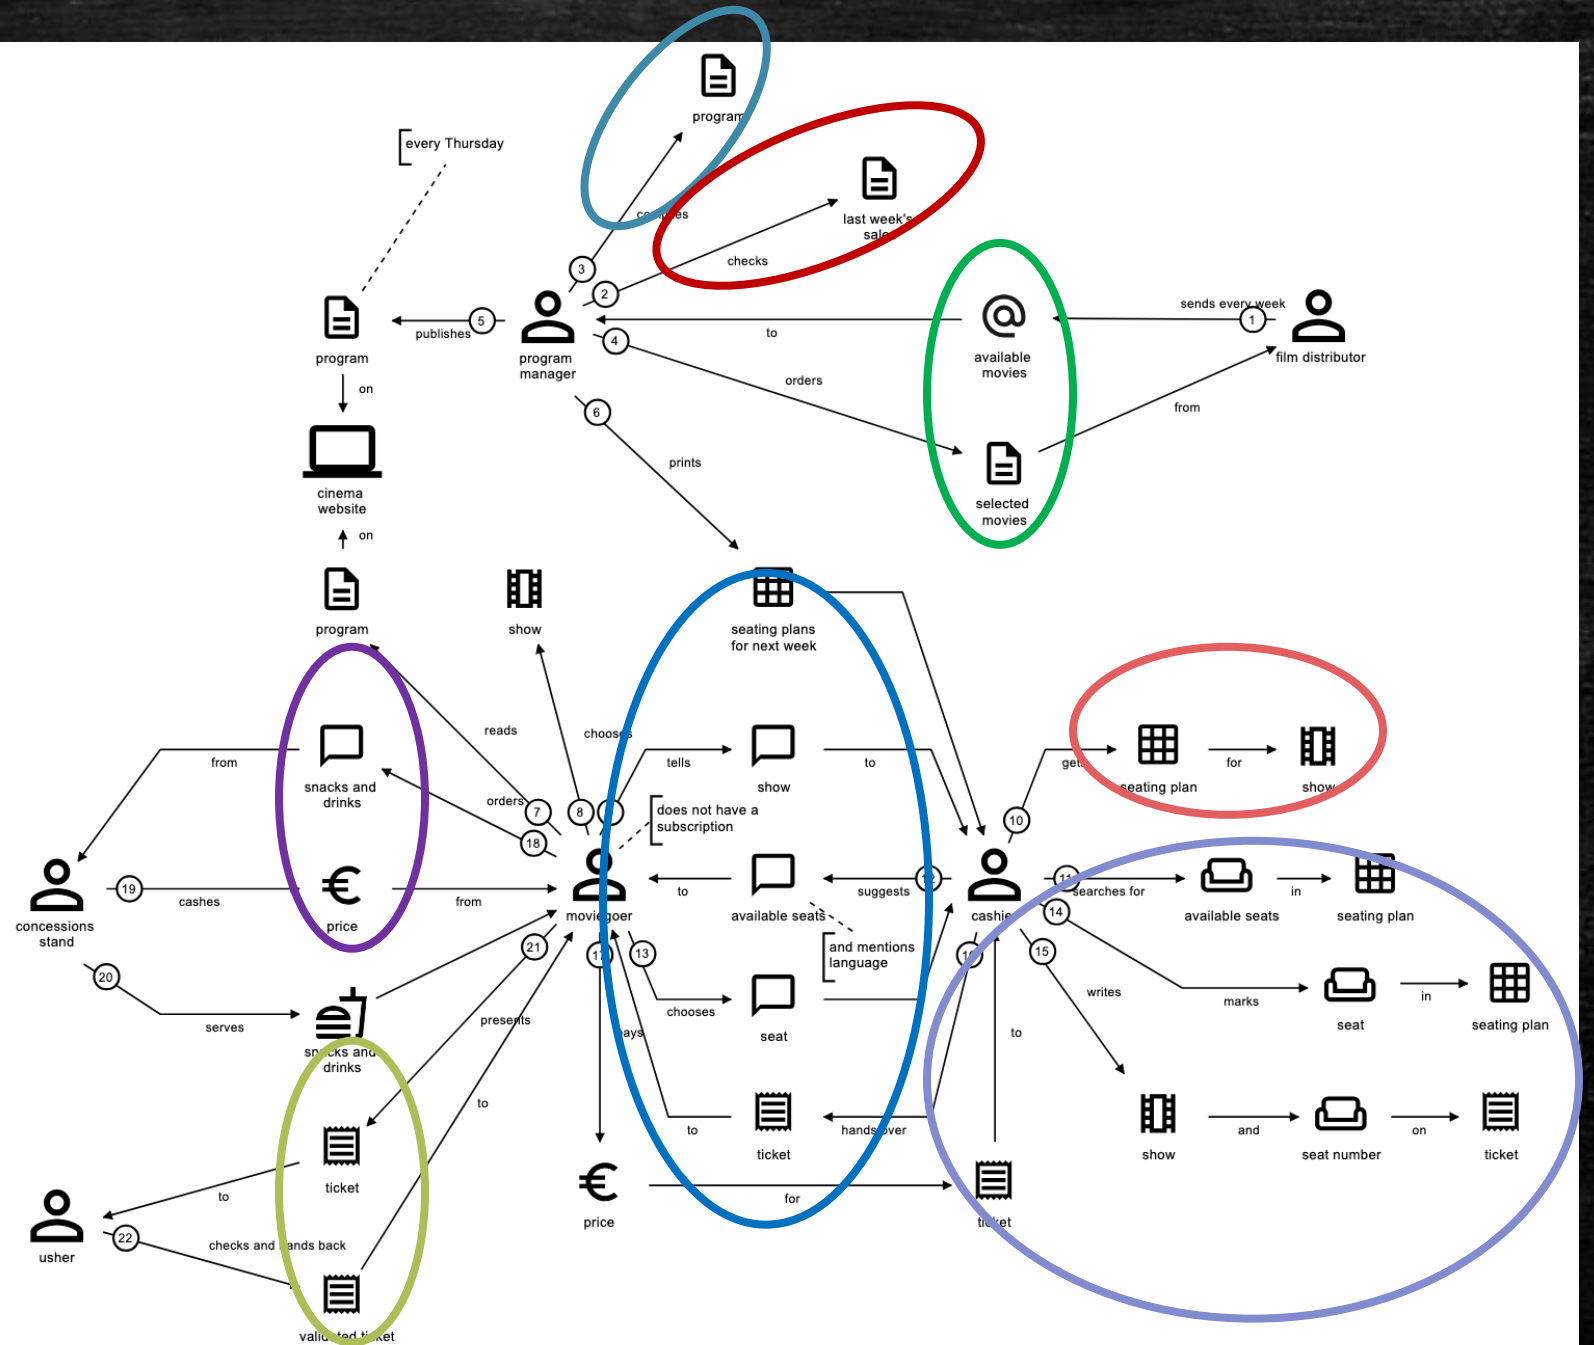
Domain Storytelling and Boundaries

Subdomains consist of activities that belong together from an actor's perspective.

- Actor produces result on their own
- One-way information flow
- Different triggers
- Activities supporting something not in the picture
- Differences in language
- Different use of the same thing

Domain Storytelling

Examples




Logical

The Domain will change as the business changes:

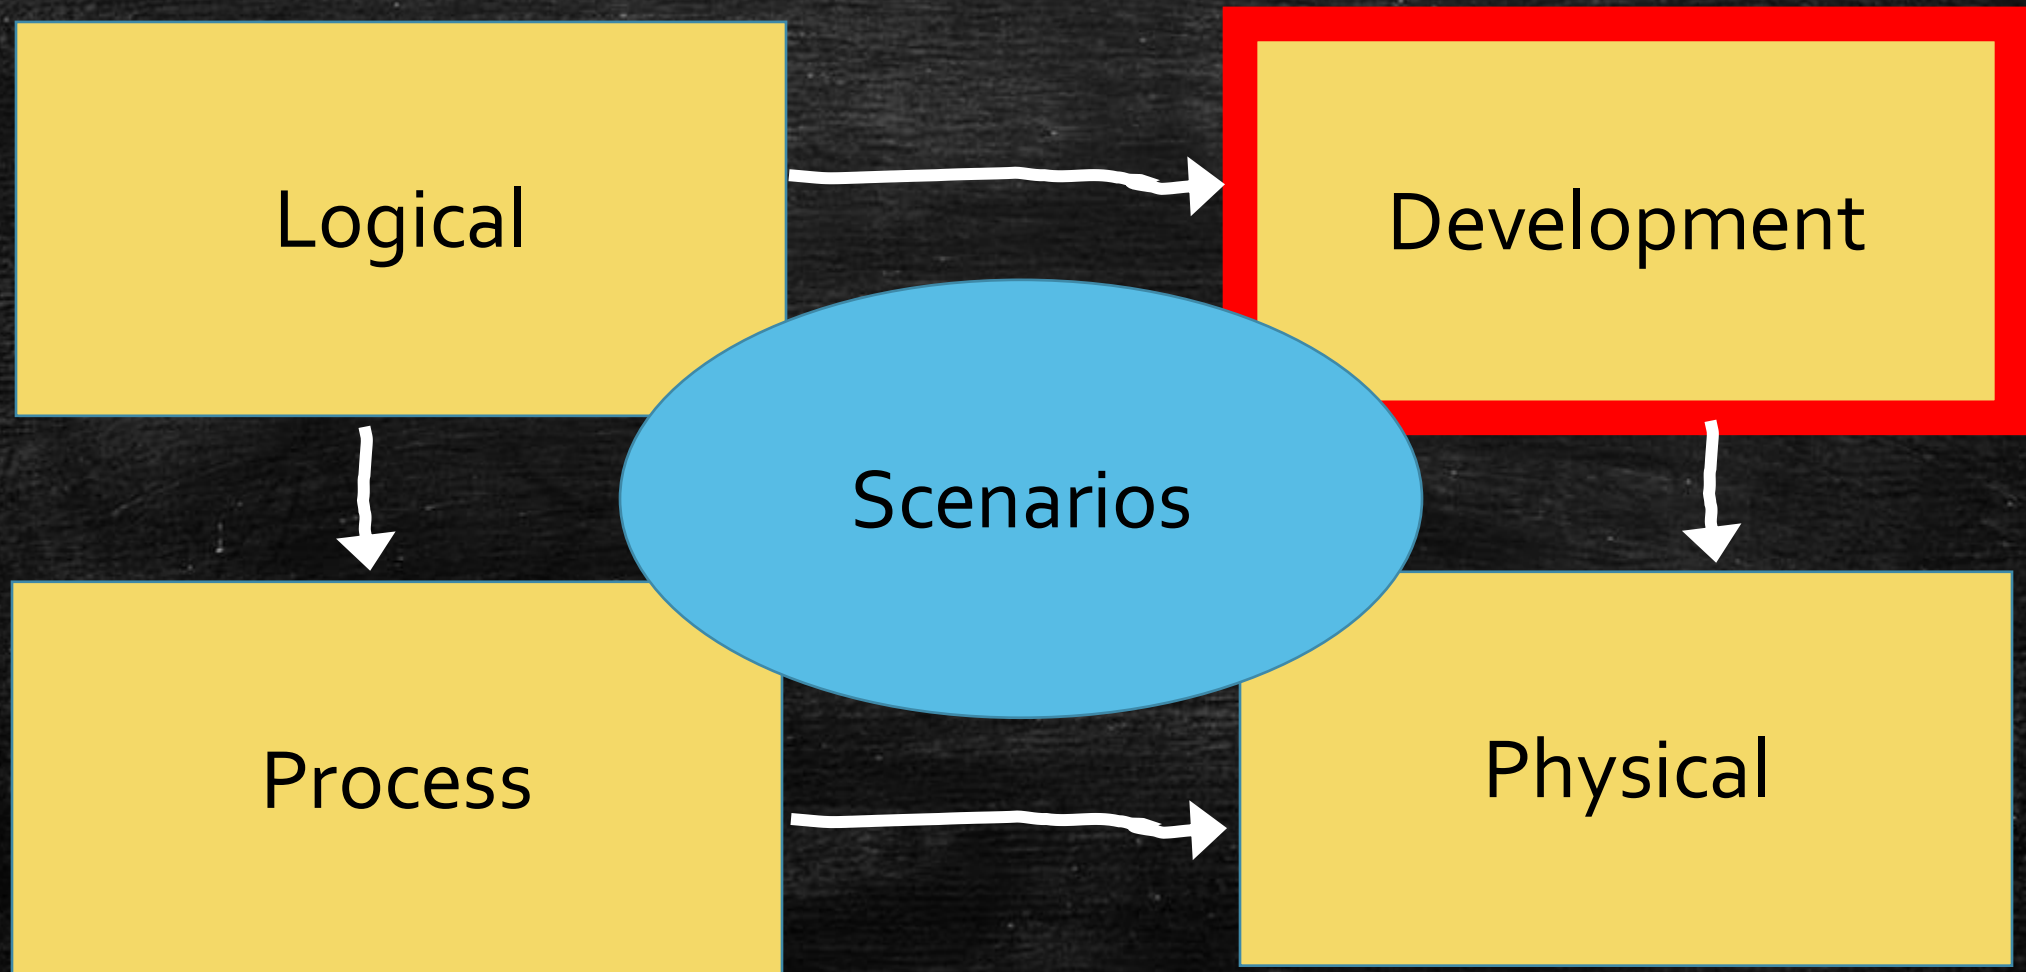
- The changes to our domain will be a constant force; a successful business is not static
 - Tasks and activities change more frequently, are retired, replaced or introduced
 - Processes change more slowly, are less frequently retired, replaced; may be more frequently introduced as new opportunities emerge.
- Our logical view requires constant gardening
 - We need a continuous conversation to share the logical model as it changes.

YOU MUST BE THIS TALL
TO USE MICROSERVICES.



The “4 +1 “ View Model - Philippe Kruchten

t: 25



- See Chris Simon, [Experiences scaling a modular monolith to microservices using the 4+1 views](#)

Development

The development architecture focuses on the actual software module organization on the software development environment. The software is packaged in small chunks—program libraries, or subsystems — that can be developed by one or a small number of developers.

Kruchten, Philippe Architectural Blueprints the 4 +1 View Model of Software


This is where we define the service boundaries that map to the logical model

Many designers have made attempts to chop a system into manageably small pieces; unfortunately, they have often found that implementation time increased rather than decreased. The key frequently lies in the second part of our stipulation above: **The parts of the original problem must be solvable separately.**

Yourdon, Edward; Constantine, Larry Structured Design, 1975

- Intramodule effects
- Intermodule effects
- Total Cost

Yourdon, Edward; Constantine, Larry
Structured Design 1975



Granularity Integrators and Disintegrators

Granularity Disintegrators

Service scope too large

Distribution of code changes

Continuous change for new problems

Granularity Integrators

ACID transactions

Shared Code

Shared Data

After Ford, Neal; Richards, Mark; Sadalage, Pramod; Dehghani, Zhamak. Software Architecture: The Hard Parts

Development

A fracture plane is a natural seam in the software system that allows the system to be split easily into two or more parts.... It is usually best to try to align software boundaries with the different business domain areas...

Fracture Plane: Business Domain Bounded Context

Skelton, Matthew; Pais, Manuel. Team Topologies

Multiple [domain] models are in play on any large project. Explicitly define the context within which a [domain] model applies. **Explicitly set boundaries in terms of team organization, usage within specific parts of the application, and physical manifestations such as code bases and database schemas.** Keep the [domain] model strictly consistent within these bounds, but don't be distracted or confused by issues outside. A **BOUNDED CONTEXT** delimits the applicability of a particular [domain] model.

Evans, Eric Domain-Driven Design

Development

First model: multiple bounded contexts, one process

Use a monolith

- It may be that you are not a large enough org, or too immature to use microservices
 - a monolith; use components within the monolith for domain boundaries
 - the system has a single CI boundary
 - Typically, has no more than 35 devs/5 teams working on the monolith (10 pizza rule?)
 - Will need a monorepo.
- If you are unsure of the bounded contexts in your domain, easier to discover here
 - Boundaries within a process are more supple; easier to change; easier to violate.

Development

Simple model: one bounded context per subdomain (two-pizza team).

Use a microservice

- Code for the subdomain fits neatly into your head
- Forces for change or replacement are cohesive: keep what changes together, together
- ACID requirements can be satisfied within the service boundary
- You have sufficient scale to justify componentization via services
 - The microservice approach to division is different, splitting up into services organized around business capability.

Development

Complex domain model: domain too large for two pizza team

- Is the code for the domain larger than your head?
 - Can you find seam(s) within the domain and split into multiple bounded contexts
 - Within a business process look for activities within the process (in domain storytelling deal with different work items)
 - Return to microservices but beware complexity of intermodule dependencies
- Would ACID transactions risk propagating across a service boundary if you split?
 - You can do BASE across a boundary but not ACID; suggests a single service
- Does the domain require a single bounded contexts (not independently solvable)?
 - Includes shared kernel model of bounded context dependency

Development

Complex model: multiple small domains, one two-pizza team.

- Code for each domains is much smaller than your head
- May need to be combined with other domains in one bounded context
 - Domain is not truly independent of others. Does not have separate reasons to change.
 - The behaviours are about the same work item
- We have tight chatty communication
- We would distribute an ACID transaction
- We would weaken cohesion as a separate context

Development

Use a macroservice

- a service that has components within it (that don't use componentization via services)
 - Not a monolith - the system has more than one service
 - Has no more than 20 devs/3 teams working on any macroservice (5 pizza rule?)
 - May or may not have/need monorepo.
 - Dependency management becomes a lot easier (though still non-trivial) the fewer the services/repos
 - better observability, debugging

Sridharan, Cindy [@copyconstruct](#) [X](#)

Development

We identified that our domain will change. This implies that the bounded contexts we have drawn will move out-of-alignment

- Our development view requires constant gardening


When your components are services with remote communications, then refactoring is much harder than with in-process libraries.

Lewis, James Fowler ,Martin - [Microservices](#)

A hyperliminal system is a system where a complicated, ergodic, ordered system executes inside a complex, non-ergodic, disordered context... It is the job of the architect to design a structure for the complicated system that will allow it to survive as the complex context around it changes and moves.

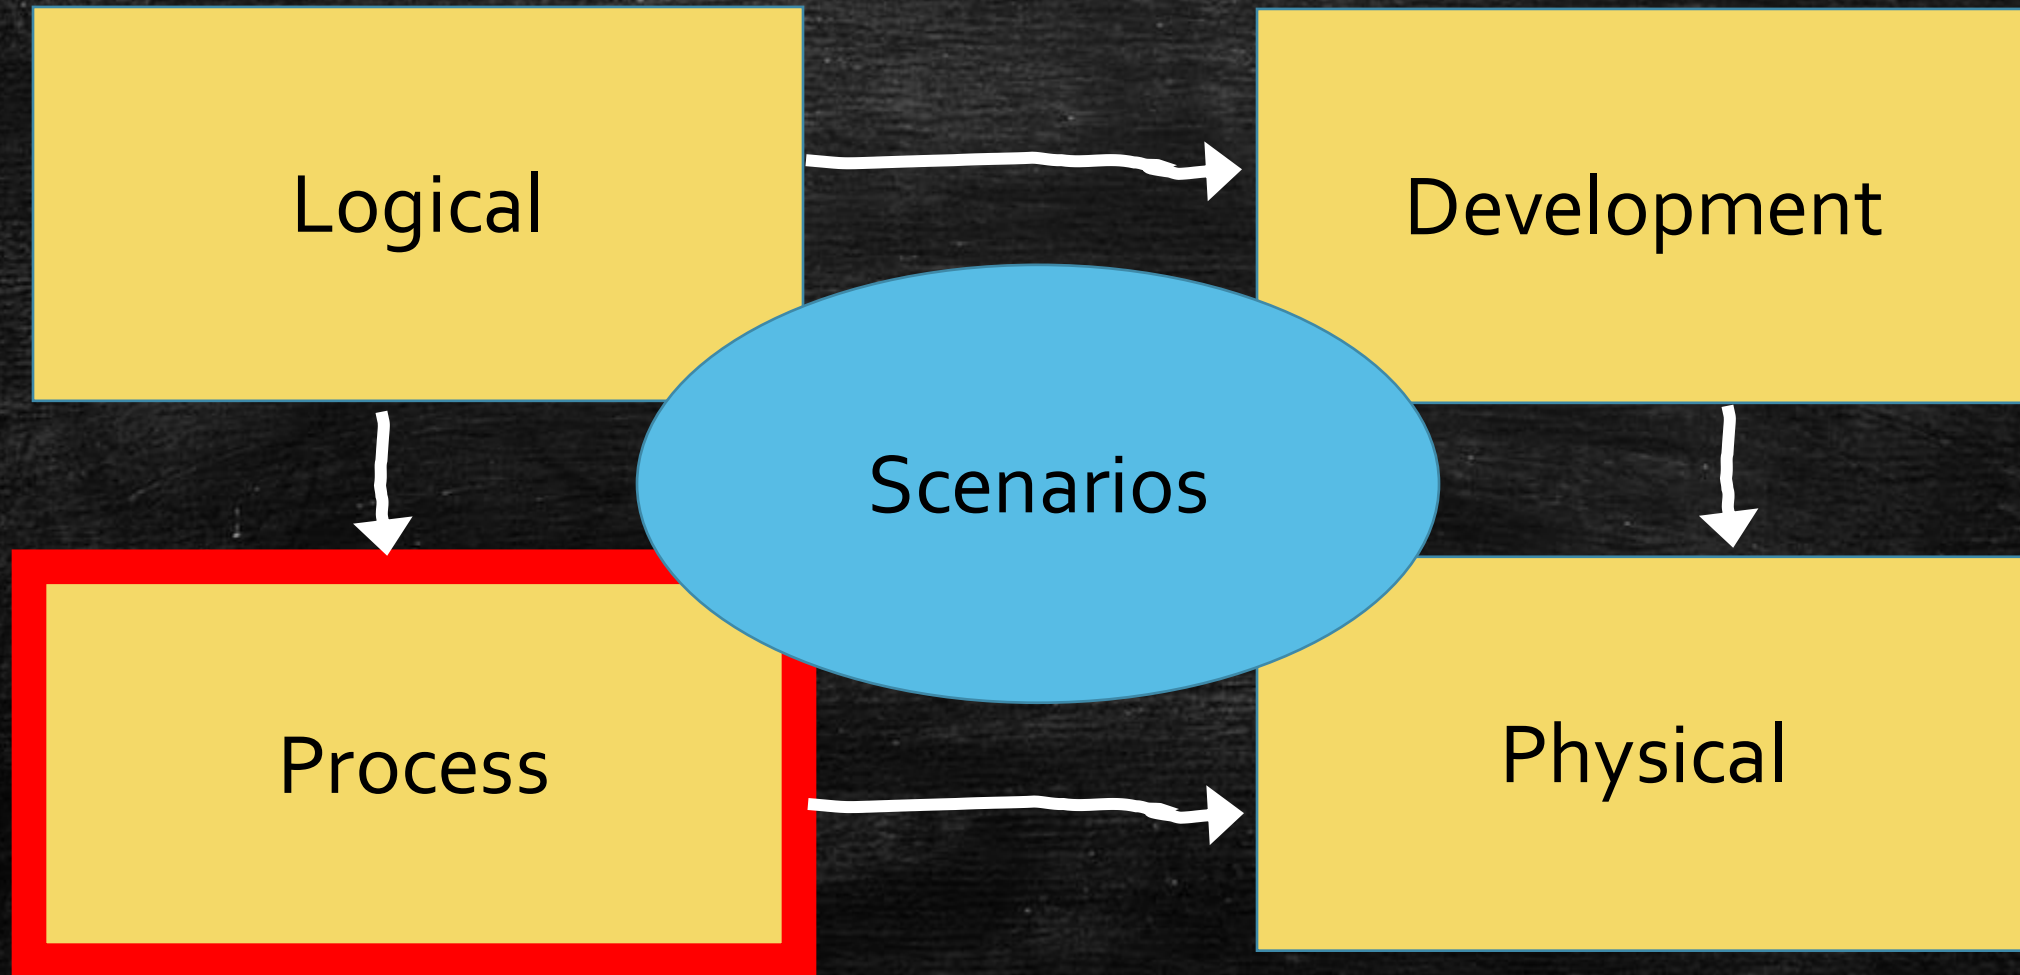
O'Reilly Barry - Residues

YOU MUST BE THIS TALL TO
USE MICROSERVICES.



The "4 +1 " View Model - Philippe Kruchten

t: 40



- See Chris Simon, [Experiences scaling a modular monolith to microservices using the 4+1 views](#)

Process

The process architecture takes into account some non-functional requirements, such as performance and availability. It addresses issues of concurrency and distribution, of system's integrity, of fault-tolerance, and how the main abstractions from the logical view fit within the process architecture...**A process is a grouping of tasks that form an executable unit. Processes represent the level at which the process architecture can be tactically controlled (i.e., started, recovered, reconfigured, and shut down)**

Kruchten, Philippe Architectural Blueprints the 4 +1 View Model of Software

This is where we define the processes that make up a microservice

Process

This is about how our processes interoperate with each other and infrastructure

- Seek to be application-centric, not infrastructure centric
- 12-Factor Apps

Interoperability

- Asynchronous communication
 - Messaging: Change or query things
 - Eventing: Distribute facts
- Synchronous communication
 - REST: Expose resources via uniform interface
 - RPC: Expose functions remotely

Process


At a first approximation, we can observe that services map to runtime processes, but that is only a first approximation. A service may consist of multiple processes that will always be developed and deployed together, such as an application process and a database that's only used by that service.

Lewis, James Fowler ,Martin - [Microservices](#)

Process

In the twelve-factor app, processes are a first class citizen. Processes in the twelve-factor app take strong cues from the unix process model for running service daemons.

Using this model, the developer can **architect their app to handle diverse workloads** by assigning **each type of work** to a **process type**. For example, HTTP requests may be handled by a web process, and long-running background tasks handled by a worker process... **the array of process types and number of processes of each type is known as the process formation**



Granularity Integrators and Disintegrators

Granularity Disintegrators

Scalability and throughput

Fault tolerance

Granularity Integrators

ACID transactions


Shared Code

After Ford, Neal; Richards, Mark; Sadalage, Pramod; Dehghani, Zhamak. Software Architecture: The Hard Parts

Process

μ

=





Microservice


App

M

=



Or



Plugin

Macroservice


App

App

Microkernel



Or



Plugin



=



Process Formation

App


Microkernel

Process
Type

Process
Type

Process
Type

YOU MUST BE THIS TALL TO USE
MICROSERVICES




3.0

5.0


5.0

5.0



The "4 +1 " View Model - Philippe Kruchten

t: 55s



- See Chris Simon, [Experiences scaling a modular monolith to microservices using the 4+1 views](#)

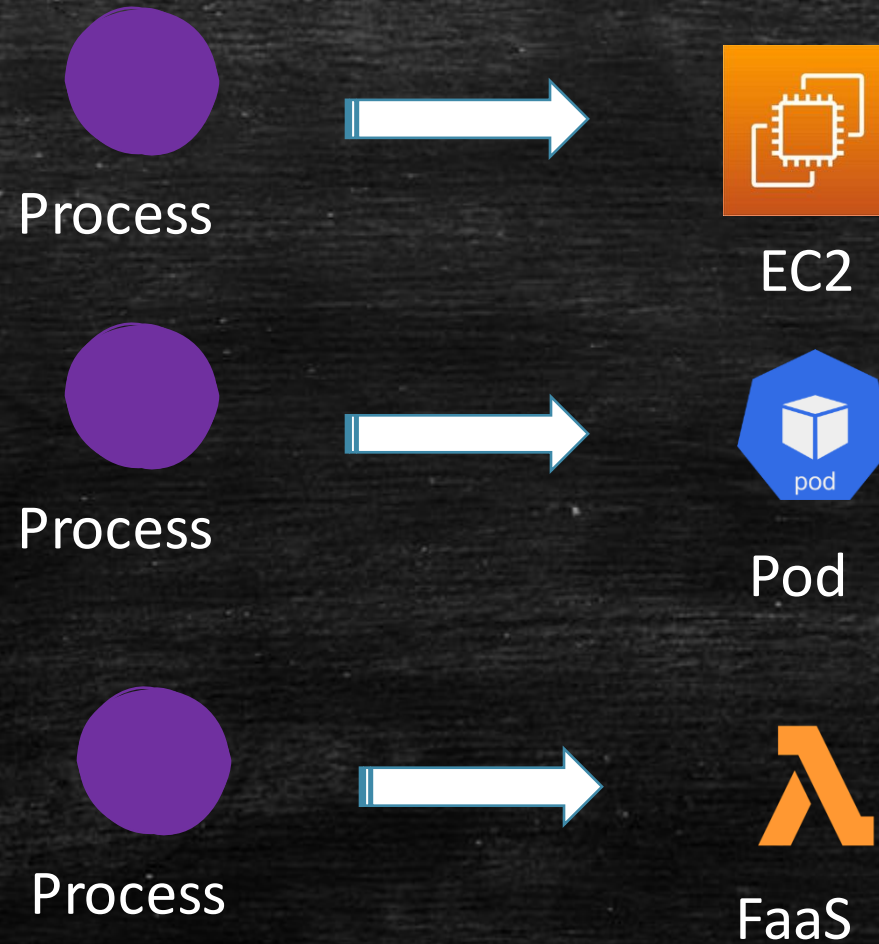
Physical

The software executes on a network of computers, or processing nodes (or just nodes for short). The various elements identified—networks, processes, tasks, and objects—need to be mapped onto the various nodes. We expect that several different physical configurations will be used: some for development and testing, others for the deployment of the system for various sites or for different customers.

Kruchten, Philippe Architectural Blueprints the 4 +1 View Model of Software

This is where we map the processes to hosts

Physical



Q&A

“Entities are not to be multiplied beyond necessity.” William of Ockham